

## THEORY ASSIGNMENT 2

ABHAY SHANKAR K: CS21BTECH11001

(1) 6.8

```
1          void bid(double amount) {
2              if (amount > highestBid)
3                  highestBid = amount;
4          }
```

A race condition can be caused when two `bid()` calls and let occur simultaneously, and both `amount` in both cases is greater than `highest bid`. Let these calls be  $c_1$  and  $c_2$ ,  $c_1.amount > c_2.amount$ . So if  $c_1$  enters the `if` block in 2 first, and then  $c_2$  enters. Then  $c_1$  executes 3, setting `highestBid` to the correct value,  $c_1.amount$ . Lastly,  $c_2$  executes 3 (which it ought to, since it has already entered the `if` block), setting the value of `highestBid` to  $c_2.amount$ .

So, race conditions cause problems.

(2) 6.11

Yes, the given ‘compare and CAS’ idiom is viable. Since any number of threads can read a shared variable, the surrounding `if` block does not impact the efficacy of the algorithm. Since the `compare_and_swap()` function is still called within the `if` block, no two threads can acquire the lock simultaneously. The extra branch instruction will preserve mutual exclusion.

(3) 6.23

An implementation of a limited open-socket-count with semaphores is fairly straightforward. We can create a semaphore initialised to  $N$ , and the code for accepting a new connection can be placed between a `sem_wait()` call and a `sem_post()` call. If more requests arrive, they will be blocked by the wait call.

(4) Reader - writer.

```
1  semaphore rw_mutex = 1;
2  semaphore mutex = 1;
3  int read_count = 0;
4  int history = 0;
5
6  writer() {
7      while (true) {
8          wait(rw_mutex);
9          history = 0;
10         ...
11         /* writing is performed */
```

```

12          ...
13          signal(rw_mutex);
14      }
15  }
16
17  reader() {
18      while (true) {
19          while(history > 20); // keeps incoming threads occupied,
                                // ensures that semaphore released in 34 is not
                                // taken up by a reader > 20 times in a row.
20          wait(mutex);
21          read_count++;
22          history++;
23          if (read_count == 1 || history == 1)
24              wait(rw_mutex);
25          signal(mutex);
26
27          ...
28          /* reading is performed */
29          ...
30
31          wait(mutex);
32          read_count--;
33          if (read_count == 0 || history > 20) // too many readers
                                                // executed in a row,
                                                // gives writer a chance.
34              signal(rw_mutex);
35          signal(mutex);
36      }
37  }

```

If twenty readers acquire the lock in a row, then further requests by readers are blocked (busy wait), and once all active readers terminate, a writer (since no readers are waiting on the lock) will enter. It then resets the **history** variable, allowing readers to make requests for the lock again. Once the writer finishes, any one of the reader/writer threads may acquire the lock.

- (5) Atomic increment - using critical section. Starvation-freedom ensured through bounded waiting.

```

1  int incr(int n) {
2      while (true) {
3          waiting[i] = true;
4          key = 1;
5          while (waiting[i] && key == 1)
6              key = compare_and_swap(&lock, 0, 1);

```

```
7      waiting[i] = false;
8      n++;
9      j = (i + 1) % n;
10     while ((j != i) && !waiting[j])
11         j = (j + 1) % n;
12         if (j == i)
13             lock = 0;
14         else
15             waiting[j] = false;
16     }
17 }
```